

Generative Adversarial Nets

Ian J. Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley,
Sherjil Ozair, Aaron Courville, Yoshua Bengio

Presenters

Do Quang Phong - 15021797

Nguyen Tuan Anh - 15020971

Nguyen Hoai Nam - 18020929

Nguyen Minh Dat - 18020288



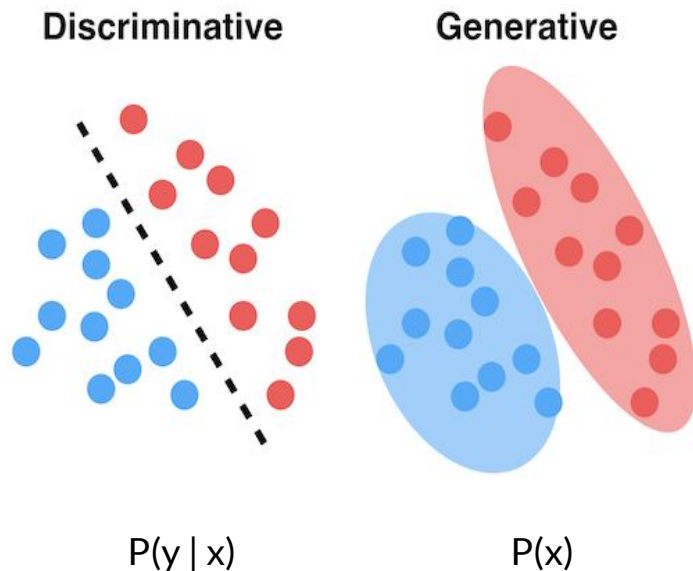
Outline

- Introduction
- Generative Adversarial Nets
- Advancements
- Applications
- Conclusion

Introduction

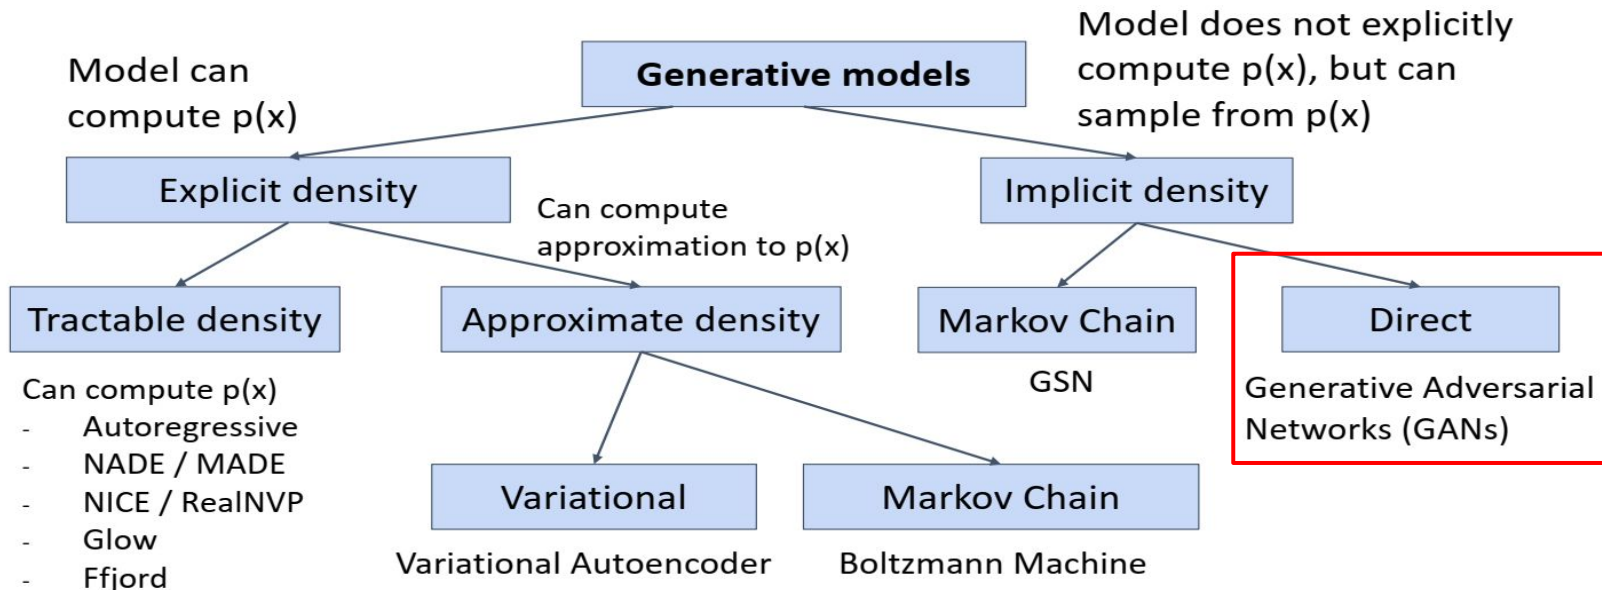
Generative vs Discriminative Models

- Given a data sample x , Discriminative model aims at predicting its label y , hence its model by the posterior distribution $P(y|x)$
- Generative models instead model the distribution $P(x)$ defined over the datapoints x .



Generative Adversarial Networks

Taxonomy of Generative Models



GANs Over Time



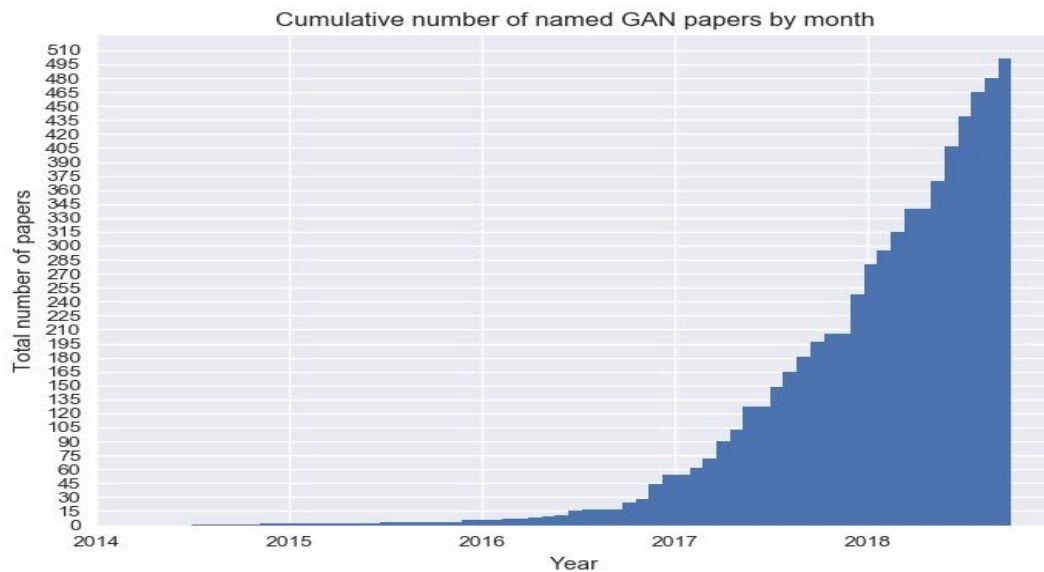
2014 - Black - White



2018 - Colored
Photo Realistic

4.5 years of GAN progress on face generation

GANs Over Time




Generative Adversarial Networks (GANs)

- An introduction to Generative Adversarial Networks (GANs Goodfellow) accepted in NIPS 2014.
 - Generator, discriminator.
 - Some GAN variants



Ian Goodfellow



Generative Adversarial Networks (GANs)


Generator

Generator learns to make **fakes** that look like **real**



Discriminator

Discriminator learns to distinguish **real** from **fake**

Fake



Real

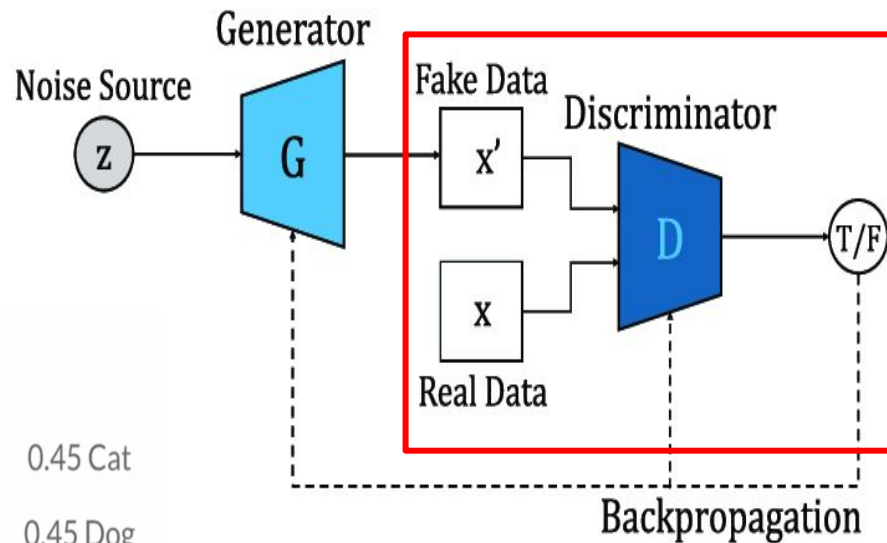



Fake



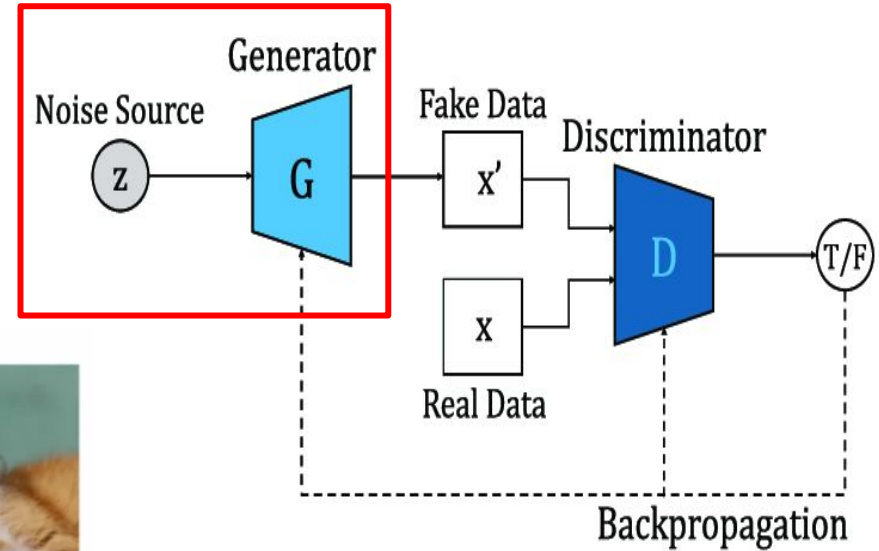

Discriminator

- The discriminator is a classification system
- It learns the probability of class Y (real or fake) given by features X.
- The probabilities are the feedback for generator



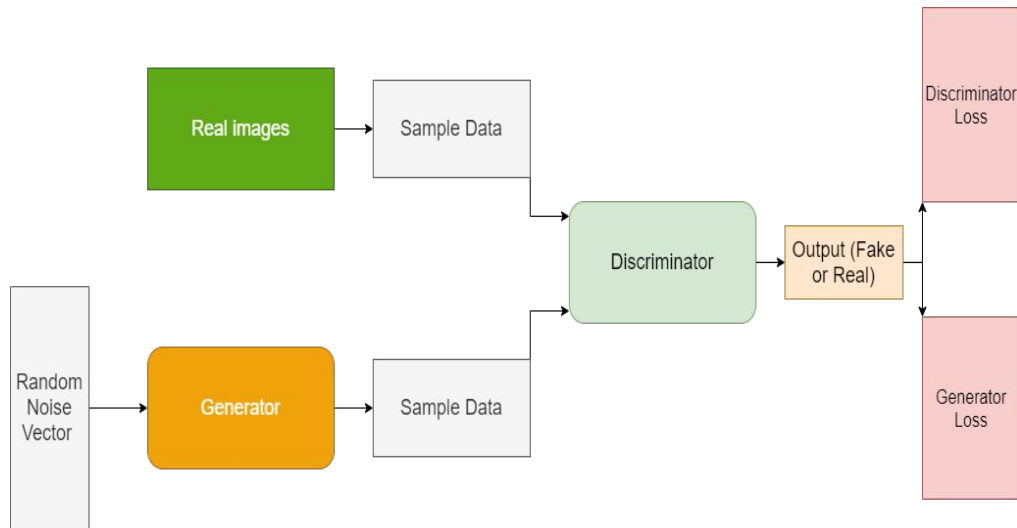
Generator

- The generator produces fake data
- It learns about the probability of features X
- The generator takes input as noise (random features)



Putting It All Together


- The Binary Cross Entropy Loss is used for training model.
- GANs training in an alternating fashion
- Train Generator & Discriminator together, like a minimax game (adversarial training)



GAN Training Phase

In each iteration, alternately train

- Discriminator
 - Learn a real sample : x
 - Learn a fake sample : $G(z)$
- Generator
 - Optimize G to fool D
 - $G(z) \rightarrow \text{Real}$
- Object Function



$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))].$$

Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k , is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Sample minibatch of m examples $\{x^{(1)}, \dots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(x^{(i)}) + \log \left(1 - D(G(z^{(i)})) \right) \right].$$

end for

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Update the generator by descending its stochastic gradient:

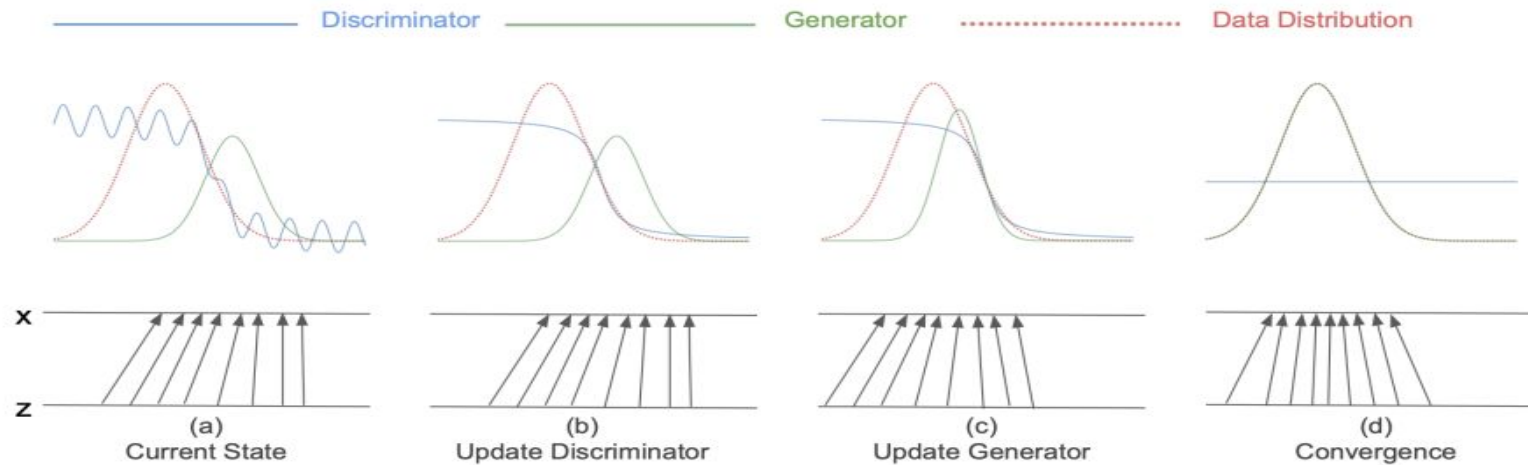
$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log \left(1 - D(G(z^{(i)})) \right).$$

end for

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

GAN Training Phase

Alternating optimization in GANs





Theoretical Analysis

- For a given generator, the optimal discriminator is:

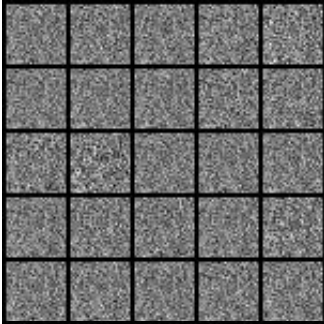
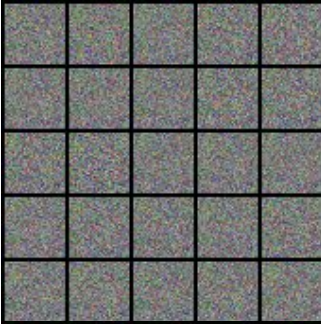
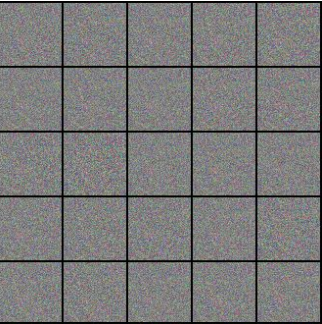
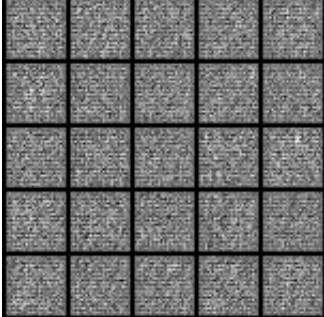
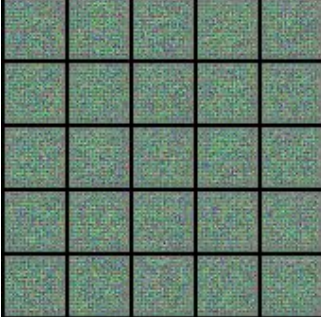
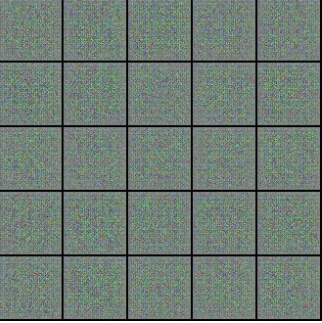
$$D_G^*(\mathbf{x}) = \frac{p_{data}(\mathbf{x})}{p_{data}(\mathbf{x}) + p_g(\mathbf{x})} \quad (2)$$

Proof. The training criterion for the discriminator D , given any generator G , is to maximize the quantity $V(G, D)$

$$\begin{aligned} V(G, D) &= \int_{\mathbf{x}} p_{data}(\mathbf{x}) \log(D(\mathbf{x})) dx + \int_{\mathbf{z}} p_{\mathbf{z}}(\mathbf{z}) \log(1 - D(g(\mathbf{z}))) dz \\ &= \int_{\mathbf{x}} p_{data}(\mathbf{x}) \log(D(\mathbf{x})) + p_g(\mathbf{x}) \log(1 - D(\mathbf{x})) dx \end{aligned} \quad (3)$$

For any $(a, b) \in \mathbb{R}^2 \setminus \{0, 0\}$, the function $y \rightarrow a \log(y) + b \log(1 - y)$ achieves its maximum in $[0, 1]$ at $\frac{a}{a+b}$. The discriminator does not need to be defined outside of $Supp(p_{data}) \cup Supp(p_g)$, concluding the proof. \square

Experiments


	MNIST	CIFAR10	CELEBA
Vanilla GAN			
DCGAN			

Advancements

Conditional GAN (cGAN)

- Generate image x with known label y
- Input noise vector with one-hot presentation of y to both G & D
- Discriminator loss by class

Result on MNIST






Deep Convolutional GAN (DCGAN)

- First working CNN-based GAN
- Notably, the generator uses “*transposed convolutional layers*” also informally called “*Deconvolution layers*”

Architecture guidelines for stable Deep Convolutional GANs


- Replace any pooling layers with strided convolutions (discriminator) and fractional-strided convolutions (generator).
- Use batchnorm in both the generator and the discriminator.
- Remove fully connected hidden layers for deeper architectures.
- Use ReLU activation in generator for all layers except for the output, which uses Tanh.
- Use LeakyReLU activation in the discriminator for all layers.

Deep Convolutional GAN (DCGAN)




Deep Convolutional GAN (DCGAN)

Original GAN (CIFAR-10)




No convolution



One convolutional layer

DCGAN (ImageNet)



Many convolutional layers
(Radford et al, 2015)



Wasserstein GAN (WGAN)

GAN Issues:

- In practice, can optimize the discriminator easier than the generator
- If distribution q far away from the ground truth p , the discriminator can saturate early & the generator barely learns anything.
- Research has indicated that if your discriminator is too good, then generator training can fail due to vanishing gradients.

Solution : Change loss function => More stable training

	Discriminator/Critic	Generator
GAN	$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right]$	$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (D(G(z^{(i)})))$
WGAN	$\nabla_w \frac{1}{m} \sum_{i=1}^m [D(x^{(i)}) - D(G(z^{(i)}))]$	$\nabla_{\theta} \frac{1}{m} \sum_{i=1}^m D(G(z^{(i)}))$

Applications



Applications


- Huge opportunity to work in applications
- Many companies are using GANs for their work
- Some applications such as :
 - Data Augmentation
 - Image Translation
 - Adversarial Machine Learning
 - Face Reenactment



Data Augmentation

- Training data issues :
 - Not enough data to achieve the best performance
 - Too imbalanced
- Supplement data when real data is ...
 - Too expensive
 - Too rare

Data Augmentation



Can mix the data augmentation techniques !




Image Translation

CycleGAN

- The most common unpaired training image-to-image translation
- Two similar generators
 - $G: x \rightarrow y$ (forward)
 - $F: y \rightarrow x$ (backward)
- Two similar discriminators D_X and D_Y
- Cycle-consistency loss




Image Translation


From one domain to another

CycleGAN



Image Translation

GauGAN



Face Reenactment




Image

+



Face landmark



Adversarial Research Areas



x
“panda”
57.7% confidence

+ .007 ×



$\text{sign}(\nabla_x J(\theta, x, y))$
“nematode”
8.2% confidence

=



$x + \epsilon \text{sign}(\nabla_x J(\theta, x, y))$
“gibbon”
99.3 % confidence

Adversarial examples



Companies using GANs



Next - gen
Photoshop



Image Filters



Text
Generation



Data
Augmentation



Conclusion

- GANs have been successfully applied to several domain and tasks
- However, working with GANs can be very challenging in practice
 - Unstable optimization
 - Mode collapse
 - Evaluation
- Besides, many cool applications with GANs are using in real life.



Acknowledgement

- [I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, Y. Bengio, “Generative Adversarial Networks”. In NIPS 2014.](#)
- [A. Radford, L. Metz, and S. Chintala. “Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks”. In ICLR 2016.](#)
- [M. Arjovsky, S. Chintala, and L. Bottou. “Wasserstein generative adversarial networks”. In ICML 2017](#)
- [CVPR 2018 Tutorials on GANs](#)
- [Generative Adversarial Networks \(GANs\) Specialization](#)
- [<https://github.com/hindupuravinash/the-gan-zoo>](#)



Further Reading

- <https://github.com/goodfeli/adversarial>
- <http://nvidia-research-mingyuliu.com/gaugan/>
- [Improved Training of Wasserstein GANs](#)
- <https://jonathan-hui.medium.com/gan-wasserstein-gan-wgan-gp-6a1a2aa1b490>
- <https://github.com/junyanz/pytorch-CycleGAN-and-pix2pix>
- [CS236G Generative Adversarial Networks](#)
- <https://blog.paperspace.com/nvidia-gaugan-introduction/>



Thank you for your attention